Natural Language Processing Assignment

2211CS020364   Narasamsetty Navadeep

AIML-THETA

1.Correct the Search Query

Explanation: Here is a basic implementation using Python, focusing on spell correction using `edit distance` and a predefined corpus of words. This code uses `zlib` for compression and `pickle` for serialization, suitable for building an offline model.

Source Code:

```python
import re

import pickle

import zlib

from collections import Counter


# Build corpus from a sample dictionary (you can enhance it with more words)

words = """going to china who was the first president of india winner of the match food in america"""


def words_list(text):

    return re.findall(r'\w+', text.lower())


WORDS = Counter(words_list(words))
# Compression for large wordlist
with open('compressed_dict. pkl', 'wb') as f:

    compressed = zlib.compress(pickle.dumps(WORDS))

    f.write(compressed)
# Load dictionary in memory
```

```python
def load_dictionary():

    with open('compressed_dict.pkl', 'rb') as f:

        return pickle.loads(zlib.decompress(f.read()))

def edit_distance_one(word):

    letters = 'abcdefghijklmnopqrstuvwxyz'

    splits = [(word[:i], word[i:]) for i in range(len(word) + 1)]

    deletes = [L + R[1:] for L, R in splits if R]

    transposes = [L + R[1] + R[0] + R[2:] for L, R in splits if len(R) > 1]

    replaces = [L + c + R[1:] for L, R in splits if R for c in letters]

    inserts = [L + c + R for L, R in splits for c in letters]

    return set(deletes + transposes + replaces + inserts)

def known(words, dictionary):

    return set(w for w in words if w in dictionary)

def candidates(word, dictionary):

    return (known([word], dictionary) or

            known(edit_distance_one(word), dictionary) or

            [word])

def correct_word(word, dictionary):

    return max(candidates(word, dictionary), key=dictionary.get)

def correct_query(query, dictionary):

    return ' '.join(correct_word(word, dictionary) for word in query.split())

# Main correction function

if __name__ == "__main__":

    dictionary = load_dictionary()

    n = int(input())

    queries = [input().strip() for _ in range(n)]

    for query in queries:

        print(correct_query(query, dictionary))
```

2. Deterministic Url and HashTag Segmentation

Explanation: This approach aims to find the most likely and meaningful segmentation of the input strings based on the provided dictionary of words and the constraint of selecting the longest valid tokens from the left.

Source Code:

```python
import re


# Load words from words.txt into a set
with open("words.txt", "r") as file:
    dictionary = set(word.strip().lower() for word in file.readlines())



def is_number(s):
    """Check if the string is a number."""
    try:
        float(s)
        return True
    except ValueError:
        return False



def tokenize(input_string, dictionary):
    """
    Tokenize the input string using the longest match first approach.


    Args:
        input_string: The string to be tokenized.
        dictionary: A set of valid words.


    Returns:
```

```python
        A list of tokens from the input string.
    """

    length = len(input_string)
    if length == 0:
        return []


    # dp[i] stores the tokens for the substring starting from index i
    dp = [None] * (length + 1)
    dp[0] = []  # Base case: empty string has no tokens


    for i in range(1, length + 1):
        # Consider all possible ending positions for the current substring
        for j in range(i):
            left_part = input_string[j:i]
            # Check if left part is a valid word or number
            if (left_part in dictionary or is_number(left_part)) and (
                dp[j] is not None
            ):
                # If left part is valid and remaining part has a valid tokenization
                right_part_tokens = dp[j]
                right_part_tokens.append(left_part)
                # Choose the longest valid tokenization
                if len(right_part_tokens) > len(dp[i]) or dp[i] is None:
                    dp[i] = right_part_tokens


    # Return the tokenization for the entire string if it exists
    return dp[length] if dp[length] is not None else [input_string]



def main():
```

```python
    """Read input strings, tokenize them, and print the results."""
    num_test_cases = int(input())
    for _ in range(num_test_cases):
        input_string = input().strip().lower()
        # Remove www and extensions for domain names, # for hashtags
        if input_string.startswith("www."):
            input_string = input_string[4:].rsplit(".", 1)[0]
        elif input_string.startswith("#"):
            input_string = input_string[1:]


        tokens = tokenize(input_string, dictionary)
        print(f"Segmentation for Input: {' '.join(tokens)}")



if __name__ == "__main__":
    main()
```

3. Disambiguation: Mouse vs Mouse

Explanation: This code provides a basic framework for classifying the usage of the word "mouse" in a sentence. You can further improve the accuracy by:

- **Expanding the Training Data:** Use a larger and more diverse dataset of sentences.
- **Experimenting with Different Classifiers:** Try other machine learning models like Support Vector Machines (SVM) or Random Forests.
- **Using Word Embeddings:** Consider using word embeddings like Word2Vec or GloVe to capture semantic relationships between words.

Source Code:

```python
import pickle

from sklearn.feature_extraction.text import CountVectorizer

from sklearn.naive_bayes import MultinomialNB


# Training data (sample corpus)

training_sentences = [

    "The complete mouse reference genome was sequenced in 2002.",

    "Tail length varies according to the environmental temperature of the mouse "

    "during postnatal development.",

    "A mouse is an input device.",

    "Many mice have a pink tail.",

    "The mouse pointer on the screen helps in navigation.",

    "A rodent like a mouse has sharp teeth.",

    "The mouse was connected to the computer using a USB port.",

    "The house was infested with mice.",

    "Computer users often prefer a wireless mouse."

]


# Labels corresponding to the training sentences
```

```python
labels = [
    "animal",
    "animal",
    "computer-mouse",
    "animal",
    "computer-mouse",
    "animal",
    "computer-mouse",
    "animal",
    "computer-mouse"
]


# Vectorize the training sentences
vectorizer = CountVectorizer()
X_train  = vectorizer.fit_transform(training_sentences)


# Create and train the Naive Bayes classifier
classifier = MultinomialNB()
classifier.fit(X_train,  labels)


# Function to predict the type of mouse
def predict_mouse_type(sentence):
    """
    Predicts whether the 'mouse' in the sentence refers to an animal or a computer mouse.

    Args:
        sentence: The input sentence.

    Returns:
        "animal" or "computer-mouse"
    """
```

```python
    vectorized_sentence = vectorizer.transform([sentence])

    prediction = classifier.predict(vectorized_sentence)[0]

    return prediction


# Get number of test cases

num_test_cases = int(input())


# Process each test case

for _ in range(num_test_cases):

    sentence = input()

    prediction = predict_mouse_type(sentence)

    print(prediction)


# Optionally, save the trained model for later use

with open('mouse_classifier.pkl', 'wb') as f:

    pickle.dump((vectorizer, classifier), f)
```

4. Language Detection

- **Explanation:** This function loads the pre-trained model from a serialized file.
- It takes a text snippet as input, normalizes it to ASCII, and converts it into a TF-IDF vector using the loaded vectorizer.
- The function then uses the trained classifier to predict the language of the snippet based on the extracted features.

Source Code:

```python
import pickle

import unicodedata

from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.naive_bayes import MultinomialNB


def normalize_to_ascii(text):
    """Remove non-ASCII characters and normalize text."""
    return unicodedata.normalize("NFKD", text).encode("ascii", "ignore").decode("ascii")


# Step 1: Training Data
training_texts = {
    "English": [
        "The quick brown fox jumps over the lazy dog.",
        "Rip Van Winkle is a story set in the years before the American Revolutionary War.",
    ],
    "French": [
        "Le renard brun rapide saute par-dessus le chien paresseux.",
        "La revolution francaise a marque une periode importante de l'histoire.",
    ],
    "German": [
        "Der schnelle braune Fuchs springt uber den faulen Hund.",
        "Die deutsche Wiedervereinigung war ein historisches Ereignis.",
    ],
```

```python
    "Spanish": [

        "El rapido zorro marron salta sobre el perro perezoso.",

        "La Revolucion Espanola fue un momento clave en la historia."

        "Si quieres que te asciendan te tienes que poner las pilas.",

    ],

}


# Normalize training data to ASCII

labels = []

texts = []

for language, samples in training_texts.items():

    labels.extend([language] * len(samples))

    texts.extend([normalize_to_ascii(sample) for sample in samples])


# Step 2: Preprocessing and Feature Extraction

vectorizer = TfidfVectorizer(ngram_range=(2, 4), analyzer="char")

X_train = vectorizer.fit_transform(texts)


# Step 3: Train the Model

classifier = MultinomialNB()

classifier.fit(X_train, labels)


# Step 4: Serialize the Model

with open("language_model.pkl", "wb") as model_file:

    pickle.dump((vectorizer, classifier), model_file)


# Step 5: Language Detection Function

def detect_language(snippet):

    with open("language_model.pkl", "rb") as model_file:

        vectorizer, classifier = pickle.load(model_file)
```

```python
    # Normalize snippet to ASCII
    snippet = normalize_to_ascii(snippet)
    X_test = vectorizer.transform([snippet])
    prediction = classifier.predict(X_test)
    return prediction[0]


# Input Processing
if _name_ == "_main_":
    # Read multi-line input
    snippet = ""
    while True:
        try:
            line = input()
            if line.strip():
                snippet += line + " "
        except EOFError:
            break

    # Predict and Output
    detected_language = detect_language(snippet.strip())
    print(detected_language)
```

5. The Missing Apostrophes

Explanation Apostrophe Handling: The code defines a function `restore_apostrophes` that iterates through each word in the input text. It uses a combination of explicit checks for common contractions (e.g., "don't," "can't," "I've") and a regular expression to handle possessive nouns (e.g., "cat's," "dog's") to restore apostrophes where appropriate.

Source Code:

```python
import re


# Function to handle apostrophes for contractions and possessives
def restore_apostrophes(text):
    restored_text = []
    words = text.split()


    for word in words:
        lower_word = word.lower()


        # Handle contractions
        if word.lower() == "dont":
            restored_text.append("don't")
        elif word.lower() == "wont":
            restored_text.append("won't")
        elif word.lower() == "cant":
            restored_text.append("can't")
        elif word.lower() == "isnt":
            restored_text.append("isn't")
        elif word.lower() == "arent":
            restored_text.append("aren't")
        elif word.lower() == "wasnt":
            restored_text.append("wasn't")
```

```python
        elif word.lower() == "werent":
            restored_text.append("weren't")
        elif word.lower() == "hasnt":
            restored_text.append("hasn't")
        elif word.lower() == "havent":
            restored_text.append("haven't")
        elif word.lower() == "hadnt":
            restored_text.append("hadn't")
        elif word.lower() == "didnt":
            restored_text.append("didn't")
        elif word.lower() == "ive":
            restored_text.append("I've")
        elif word.lower() == "were":
            restored_text.append("we're")
        elif word.lower() == "i":
            restored_text.append("I")
        elif word.lower() == "id":
            restored_text.append("I'd")
        elif word.lower() == "ive":
            restored_text.append("I've")
        elif word.lower() == "youve":
            restored_text.append("you've")
        elif word.lower() == "hes":
            restored_text.append("he's")
        elif word.lower() == "shes":
            restored_text.append("she's")
        elif word.lower() == "its":
            restored_text.append("it's")
        elif word.lower() == "were":
            restored_text.append("we're")
```

```python
        # Handle possessives (only add 's when it makes sense)
        elif re.match(r'\w+s$', word) and lower_word not in ["its", "hers", "ours", "yours", "theirs"]:
            restored_text.append(re.sub(r"s$", "'s", word))


        # For normal words that don't need apostrophes, keep them as is
        else:
            restored_text.append(word)


    return " ".join(restored_text)


# Input
input_text = """At a news conference Thursday at the Russian manned-space facility in Baikonur,
Kazakhstan, Kornienko said "we will be missing nature, we will be missing landscapes, woods." He
admitted that on his previous trip into space in 2010 "I even asked our psychological support folks to
send me a calendar with photographs of nature, of rivers, of woods, of lakes."

Kelly was asked if hed miss his twin brother Mark, who also was an astronaut.

"Were used to this kind of thing," he said. "Ive gone longer without seeing him and it was great."

The mission wont be the longest time that a human has spent in space - four Russians spent a year or
more aboard the Soviet-built Mir space station in the 1990s.

SCI Astronaut Twins

Scott Kelly (left) was asked Thursday if hed miss his twin brother, Mark, who also was an astronaut. Were
used to this kind of thing, he said. Ive gone longer without seeing him and it was great. (NASA/Associated
Press)

"The last time we had such a long duration flight was almost 20 years and of course al{-truncated-}"""


# Restore apostrophes
output_text = restore_apostrophes(input_text)

print(output_text)
```

6. Segment the Twitter Hashtags

Explanation: Tokenization with Dynamic Programming: The `segment_hashtag` function uses dynamic programming to break down the hashtag into a sequence of words. It iterates through the hashtag, checking for valid word combinations from a given dictionary and selecting the longest possible valid sequence.

Source Code:

```python
# Define a function that segments a single hashtag into words

def segment_hashtag(hashtag, word_dict):

    n = len(hashtag)

    dp = [None] * (n + 1)

    dp[0] = []  # Base case: empty string can be segmented as an empty list


    # Iterate over the hashtag string

    for i in range(1, n + 1):

        for j in range(max(0, i - 20), i): # Limit the length of words checked

            word = hashtag[j:i]

            if word in word_dict and dp[j] is not None:

                dp[i] = dp[j] + [word]

                break


    return " ".join(dp[n]) if dp[n] is not None else hashtag


# Main function to process input and output results

def process_hashtags(num_hashtags, hashtags, word_dict):

    result = []

    for hashtag in hashtags:

        segmented = segment_hashtag(hashtag, word_dict)

        result.append(segmented)

    return result
```

```python
# Sample dictionary of common words (expand this as needed)
word_dict = {
    "we", "are", "the", "people", "mention", "your", "faves",
    "now", "playing", "walking", "dead", "follow", "me"
}


# Sample input
num_hashtags = int(input())
hashtags = [input().strip() for _ in range(num_hashtags)]


# Process the hashtags and print the result
segmented_hashtags = process_hashtags(num_hashtags, hashtags, word_dict)
for segmented in segmented_hashtags:
    print(segmented)
```

7. Expand the Acronyms

Explanation: Acronym Extraction: The code extracts acronyms and their potential expansions from a given set of text snippets by identifying uppercase words within parentheses and searching for preceding phrases. It also attempts to extract acronyms not explicitly defined in parentheses by analyzing the surrounding context.

Source Code:

```python
import re


def extract_acronyms_and_expansions(snippe ts):
    """
    Extract acronyms and their expansions from the provided snippets.
    """
    acronym_dict = {}
    for snippet in snippets:
        # Find all potential acronyms (uppercase words typically enclosed in parentheses)
        matches = re.findall(r'\((\b[A-Z]+\b)\)', snippet)


        for match in matches:
            # Extract the preceding text (potential expansion)
            preceding_text = snippet.split(f"({match})")[0].strip()


            # Look for the last meaningful phrase before the acronym
            expansion_candidates = re.split(r'[.,;:-]', preceding_text)
            if expansion_candidates:
                expansion = expansion_candidates[-1].strip()
                acronym_dict[match] = expansion


        # Additionally, handle acronyms not in parentheses but defined explicitly
        words = snippet.split()
        for i, word in enumerate(words):
```

```python
            if word.isupper() and len(word) > 1:  # Likely an acronym
                if word not in acronym_dict:
                    # Try to extract its expansion from the surrounding  context
                    if i > 0:
                        preceding_context = " ".join(words[max(0, i-5):i])
                        if preceding_context:
                            acronym_dict[word] = preceding_context
    return acronym_dict


def process_tests(acronym_dict, tests):
    """
    Process test acronyms and return their expansions.
    """
    results = []
    for test in tests:
        # Normalize the test acronym (case insensitive)
        expansion = acronym_dict.get(test.upper(),  "Not Found")
        results.append(expansion)
    return results


def main():
    # Read input
    n = int(input().strip())
    snippets = [input().strip() for _ in range(n)]
    tests = [input().strip() for _ in range(n)]

    # Extract acronyms  and expansions
    acronym_dict = extract_acronyms_and_expansions(snippets)
    # Process test queries
    results = process_tests(acronym_dict,  tests)
    # Output results
```

```
    print("\n".join(results))

if _name_ == "_main_":main()
```

8. Correct the Search Query

Explanation: Here is a basic implementation using Python, focusing on spell correction using `edit distance` and a predefined corpus of words. This code uses `zlib` for compression and `pickle` for serialization, suitable for building an offline model.

Source Code:

```python
import re

import pickle

import zlib

from collections import Counter


# Build corpus from a sample dictionary (you can enhance it with more words)

words = """going to china who was the first president of india winner of the match food in america"""


def words_list(text):

    return re.findall(r'\w+', text.lower())


WORDS = Counter(words_list(words))


# Compression for large wordlist

with open('compressed_dict.pkl', 'wb') as f:

    compressed = zlib.compress(pickle.dumps(WORDS))

    f.write(compressed)


# Load dictionary in memory

def load_dictionary():

    with open('compressed_dict.pkl', 'rb') as f:

        return pickle.loads(zlib.decompress(f.read()))


def edit_distance_one(word):

    letters = 'abcdefghijklmnopqrstuvwxyz'

    splits = [(word[:i], word[i:]) for i in range(len(word) + 1)]
```

```python
    deletes = [L + R[1:] for L, R in splits if R]

    transposes = [L + R[1] + R[0] + R[2:] for L, R in splits if len(R) > 1]

    replaces = [L + c + R[1:] for L, R in splits if R for c in letters]

    inserts = [L + c + R for L, R in splits for c in letters]

    return set(deletes + transposes + replaces + inserts)


def known(words, dictionary):

    return set(w for w in words if w in dictionary)


def candidates(word, dictionary):

    return (known([word], dictionary) or

            known(edit_distance_one(word), dictionary) or

            [word])


def correct_word(word, dictionary):

    return max(candidates(word, dictionary), key=dictionary.get)


def correct_query(query, dictionary):

    return ' '.join(correct_word(word, dictionary) for word in query.split())


# Main correction function
if __name__ == "__main__":

    dictionary = load_dictionary()

    n = int(input())

    queries = [input().strip() for _ in range(n)]


    for query in queries:

        print(correct_query(query, dictionary))
```

G.A Text-Processing Warmup

Explanation: Article and Date Counting: The code defines a function `count_articles_and_dates` that takes a text fragment as input. It first normalizes the text to lowercase for case-insensitive article counting. Then, it uses regular expressions to count occurrences of the definite and indefinite articles ("a," "an," "the") and identify valid dates in various formats (e.g., "DD Month YYYY," "Month DD, YYYY," etc.)

Source Code:

```
import re


def count_articles_and_dates(fragment):
    """
    Count occurrences of 'a', 'an', 'the', and valid dates in a given text fragment.
    """
    # Normalize text for article counting
    lower_fragment = fragment.lower()


    # Count articles
    a_count = len(re.findall(r'\b[a]\b', lower_fragment))
    an_count = len(re.findall(r'\b[an]\b', lower_fragment))
    the_count = len(re.findall(r'\b[the]\b', lower_fragment))


    # Identify valid dates
    date_patterns = [

r'\b\d{1,2}(?:st|nd|rd|th)?(?:\s+of)?\s+(January|February|March|April|    May|June|July|August|September|October|November|December)\s+\d{2,4}\b',   # Day Month Year

r'\b(January|February|March|April|May|June|July|August| September|October|November|December)\s+\d{1,2}(?:st|nd|rd)?,?\s+\d{2,4}\b',        # Month Day Year
        r'\b\d{1,2}/\d{1,2}/\d{2,4}\b',   # Day/Month/Year
        r'\b\d{4}-\d{2}-\d{2}\b'          # ISO format: Year-Month-Day
    ]
```

```python
    # Combine all date patterns
    date_regex = '|'.join(date_patterns)
    dates = re.findall(date_regex, fragment, re.IGNORECASE)
    date_count = len(dates)


    return a_count, an_count, the_count, date_count


def main():
    import sys
    input = sys.stdin.read


    # Read input data
    data = input().strip().split("\n")
    t = int(data[0]) # Number of test cases
    fragments = data[1:] # Remaining lines contain the fragments


    results = []
    for i in range(t):
        fragment = fragments[i].strip()
        # Count articles and dates
        a_count, an_count, the_count, date_count = count_articles_and_dates(fragment)
        results.append(f"{a_count}\n{an_count}\n{the_count}\n{ date_count}")


    # Output results
    print("\n".join(results))


if _name_ == "_main_":
    main()
```

10.Who is it?

Explanation: Pronoun Identification and Entity Matching: The code first finds all pronouns (words enclosed in double backslashes) and their positions in the text. It then cleans the text by removing the backslashes. Next, it iterates through each pronoun and searches for the closest matching entity (from a provided list) that appears before the pronoun in the text.

Source Code:

```
import re

def resolve_pronouns(text, entities):
    # Extract all pronouns and their positions
    pronoun_pattern = r'\\(\w+)\\'
    pronouns = [(match.group(1), match.start()) for match in re.finditer(pronoun_pattern, text)]


    # Clean the text by removing ** markers
    clean_text = re.sub(r'\\(\w+)\\', r'\1', text)


    # Initialize a list to store the resolved entities
    resolved = []


    # For each pronoun, find the corresponding entity
    for pronoun, pos in pronouns:
        closest_entity = None
        closest_distance = float('inf')


        # Iterate through all entities to find the best match
        for entity in entities:
            entity_pos = clean_text.rfind(entity, 0, pos) # Find the last occurrence of the entity before the pronoun
            if entity_pos != -1:
                distance = pos - (entity_pos + len(entity))
                if distance < closest_distance:
```

```python
            closest_distance = distance
            closest_entity = entity

    # Append the resolved entity to the list
    resolved.append(closest_entity)


    return resolved



def main():
    import sys
    input = sys.stdin.read
    data = input().strip().split("\n")


    # Read the number of lines in the text snippet
    n = int(data[0])


    # Combine the next N lines into the full text snippet
    text_snippet = " ".join(data[1:n + 1])


    # Read the list of entities
    entities = [e.strip() for e in data[n + 1].split(';')]


    # Resolve pronouns
    result = resolve_pronouns(text_snippet, entities)


    # Output the resolved entities
    for entity in result:
        print(entity)
```

```python
if _name_ == "_main_":
    main()
```